

Refinements to the SAE–Thompson Sampling Framework for Active Email Outreach

Ivanov Tymofii
FIT CTU

Extends: Žid, Č., Alves, R., and Kordík, P. 2025. Active Recommendation for Email Outreach Dynamics. In Proceedings of the 34th ACM International Conference on Information and Knowledge Management (CIKM '25), Seoul, Republic of Korea. <https://doi.org/10.1145/3746252.3760832>

1 Introduction

Žid et al. [2025] frame cold-start recipient selection for a new email campaign as an iterative multi-armed bandit: each arm is a recipient, and Thompson Sampling (TS) parameters are derived from a shallow autoencoder (SAE) trained on historical recipient–template opens, avoiding retraining during the active-learning phase. Several components of that model — a fixed exploration/exploitation coefficient, an unweighted historic open rate, a linear aggregation for the current-state confidence term, training on binary labels only, and a shallow autoencoder architecture — were fixed by the original authors for simplicity. We revisit each of these choices and propose concrete, mathematically grounded refinements: dynamic scheduling of the blend coefficient, a recency-weighted historic rate, a nonlinear forward-pass confidence term, a variance-based influence score, a factored score function, two ways of exposing the autoencoder to previously unused time-to-open (TTO) signal, and a deep, nonlinear autoencoder. We report experimental results for each proposal against a reproduced baseline and outline directions for further work.

2 Recap of the Original Model

Setup. We observe $X \in \{0, 1\}^{n \times m}$, where n is the number of templates, m the number of recipients, and $X_{i,j} = 1$ if recipient j opened template i . For a new template $n+1$, sent over b batches within a deadline T , the model selects at each batch a small subset of recipients so as to maximize eventual openers reached early.

Shallow autoencoder. A SAE with encoder/decoder $E, D \in \mathbb{R}^{m \times d}$ is trained as

$$\min_{E,D} \ell(X, \sigma(XB_{E,D})), \quad B_{E,D} = ED^\top - \text{diag}([E \odot D]\mathbf{1}_m),$$

with ℓ the element-wise binary cross-entropy and the diagonal constraint ruling out trivial self-prediction. We write $f_{\text{SAE}}(x) = \sigma(x^\top B_{E,D})$ for the trained forward pass and $\Sigma = \sigma(B_{E,D})$; entry $\Sigma_{i,j}$ is the probability of recipient j opening a template given that only recipient i has opened it — a measure of i 's influence on j .

Arm score. The bandit score for recipient j at time t is

$$s_j(t) = \alpha \phi_j p_j + (1 - \alpha) f_j(t), \quad (1)$$

$$\phi_j = \frac{1}{n} \sum_{i=1}^n X_{i,j}, \quad p_j = \frac{1}{m-1} \sum_{i \neq j} \Sigma_{j,i}, \quad f_j(t) = \bar{x}(t)^\top \Sigma_{:,j}, \quad (2)$$

where ϕ_j is recipient j 's historic open rate, p_j is a general influence score, $f_j(t)$ is a confidence score specific to template $n+1$ given the observed (normalized) state $\bar{x}(t)$, and $\alpha \in [0, 1]$ trades off the two. TS parameters are then $\alpha_j(t) = G s_j(t)$, $\beta_j(t) = G(1 - s_j(t))$ for a confidence modifier G .

Two orthogonal readings of Eq. (1) motivate our refinements: $\phi_j p_j$ vs. $f_j(t)$ can be read as *exploration vs. exploitation*, or equivalently as *historic vs. current* data.

3 Proposed Refinements

3.1 Dynamic α -Scheduling

A fixed α ignores that the value of exploration should shrink as the campaign progresses, echoing the broader practice of decaying exploration coefficients over time in the bandit and reinforcement-learning literature [Auer et al., 2002]. We schedule $\alpha(t)$ against two natural progress variables: $\mu = N(t)/N$, the fraction of the sending budget already used, and $\tilde{o} = o(t)/m$, the fraction of recipients who have already opened, with $o(t) = x(t)^\top \mathbf{1}$.

Send-volume-based (function of μ). A **linear** schedule from l to r , in the spirit of the linearly annealed ε -greedy exploration rate used to train DQN agents [Mnih et al., 2015],

$$\alpha(\mu) = l(1 - \mu) + r\mu, \quad (3)$$

and a **geometric** (log-linear) schedule, analogous to the geometric cooling schedules of simulated annealing [Kirkpatrick et al., 1983],

$$\alpha(\mu) = l^{1-\mu} r^\mu. \quad (4)$$

Open-rate-based (function of $\tilde{o}$). Since $f_j(t)$ is itself a point estimate that becomes more reliable as more opens are observed, we also consider a **confidence** schedule, hyperbolic in $\tilde{o}$ — paralleling how the confidence radius of UCB-style bandit policies shrinks as observations accumulate [Auer et al., 2002]:

$$\alpha(\tilde{o}) = \frac{\kappa}{\tilde{o} + \kappa}, \quad (5)$$

with dampening factor κ (so $\alpha(\kappa) = \frac{1}{2}$). All three reduce to the constant baseline at their degenerate limits ($l = r$, or $\kappa \rightarrow \infty$). A schedule jointly conditioned on both μ and $\tilde{o}$ is conceptually natural but is left to future work (Section 5) to avoid an unconstrained hyperparameter count.

3.2 Recency-Weighted Historic Rate ϕ_j

The original ϕ_j weights every past template equally, ignoring drift in recipient behavior over time — a limitation well studied in time-aware collaborative filtering [Ding and Li, 2005, Koren, 2009]. Assuming templates $X_1, \dots, X_n$ were sent at uniform intervals, we assign template i an exponential weight, following the time-weighting scheme of Ding and Li [2005],

$$\omega_i = 2^{-d_i/h}, \quad d_i = \frac{n-i}{n}, \quad (6)$$

with half-life $h > 0$, and redefine

$$\phi_j = \frac{\omega^\top X_{:j}}{\sum_i \omega_i}. \quad (7)$$

Normalizing by $\sum_i \omega_i$ keeps $\phi_j \in [0, 1]$. As $h \rightarrow \infty$ every $\omega_i \rightarrow 1$, recovering the original unweighted average as a special case.

3.3 Forward-Pass Definition of $f_j(t)$

$f_j(t) = \bar{x}(t)^\top \Sigma_{:j}$ is a linear combination of individual influences $\Sigma_{i,j}$ and therefore captures only the *average* effect of openers, never their joint interaction. We instead pass the full current state through the trained network, in the spirit of nonlinear autoencoder-based collaborative-filtering predictors such as AutoRec [Sedhain et al., 2015]:

$$f_j(t) = \sigma(x(t)^\top B_{E,D}) e_j = f_{\text{SAE}}(x(t)) e_j. \quad (8)$$

Because σ is nonlinear, $\sigma(x(t)^\top B_{E,D}) \neq x(t)^\top \sigma(B_{E,D})$ in general (equality holds only for one-hot $x(t)$), so Eq. (8) is a genuinely different, jointly-conditioned probability estimate rather than a relabeling of Eq. (2).

3.4 Variance-Based Definition of p_j

p_j is meant to reward recipients whose opening is *informative* about others, in the sense of informativeness used throughout active learning [Settles, 2009]. The original definition, $p_j = \frac{1}{m-1} \sum_{i \neq j} \Sigma_{j,i}$, rewards recipients that raise the *average level* of other recipients' f -scores — which is not the same as reducing uncertainty about them. We instead reward maximal *spread* induced across other recipients' scores, following variance-based exploration bonuses in the multi-armed bandit literature [Audibert et al., 2009]:

$$p_j = \frac{4}{m-1} \sum_{i \neq j} (\Sigma_{j,i} - \bar{\Sigma}_{j:})^2, \quad \bar{\Sigma}_{j:} = \frac{1}{m-1} \sum_{i \neq j} \Sigma_{j,i}. \quad (9)$$

This is the variance of $\Sigma_{j:}$ (excluding j), scaled by 4 so that the maximum possible variance of a Bernoulli variable (0.25) maps to 1, keeping $p_j \in [0, 1]$.

3.5 Factored Score $s_j(t) = \pi_j(t) u_j(t)$

The linear combination in Eq. (1) conflates the two orthogonal axes noted in Section 2 — it is not clear, for instance, why exploration should draw solely on historic data while exploitation draws solely on current data. We instead factor the score into a *probability-of-opening* term and a *utility-of-opening* term, each independently schedulable:

$$\pi_j(t) = \alpha(t) \phi_j + (1 - \alpha(t)) f_j(t), \quad u_j(t) = \beta(t) p_j + (1 - \beta(t)) \cdot 1, \quad (10)$$

$$s_j(t) = \pi_j(t) u_j(t) = \left(\alpha(t) \phi_j + [1 - \alpha(t)] f_j(t) \right) \left(\beta(t) p_j + [1 - \beta(t)] \right). \quad (11)$$

Here $\pi_j(t)$ blends historic and current opening-probability estimates, while $u_j(t)$ blends the direct value of an open (the constant 1) with its indirect information value p_j . This factorization reuses every primitive introduced above unchanged, so refinements from Sections 3.1–3.4 can be dropped into $\pi_j(t)$ and $u_j(t)$ directly — we use a confidence schedule (Eq. (5)) for $\alpha(t)$ and a send-volume schedule (Eq. (3)/(4)) for $\beta(t)$.

3.6 Incorporating Time-to-Open (TTO)

The SAE is trained purely on binary opens, discarding *how fast* a recipient opened — signal shown to carry predictive value for engagement [Sinha et al., 2018] and that matters given the short operational window $(T/b, T)$. Let $\Delta_{i,j} \in \mathbb{R}_{\geq 0} \cup \{+\infty\}$ be the observed time-to-open (with $\Delta_{i,j} = +\infty$ for non-openers, so $X_{i,j} = \mathbb{1}[\Delta_{i,j} < \infty]$), and define a decayed open label

$$Y_{i,j} = 2^{-\Delta_{i,j}/h_\delta}, \quad 2^{-\infty} := 0, \quad (12)$$

with TTO half-life h_δ ; as $h_\delta \rightarrow \infty$, $Y_{i,j} \rightarrow X_{i,j}$ for every opener.

Binary-to-TTO autoencoder. We keep the SAE’s binary input but replace the reconstruction target:

$$\min_{E,D} \ell(Y, \sigma(XB_{E,D})). \quad (13)$$

Since the input space is unchanged, the forward-pass definitions of p_j and $f_j(t)$ (Eqs. (8)–(9)) carry over mechanically; only their interpretation shifts — $\Sigma_{j,i}$ now blends *whether* i opens with *how fast*.

TTO cutoff. Training rows of X reflect the *eventual* open pattern, while active-learning queries $x(t)$ are necessarily censored at $t < T$. We instead censor the training input itself, fixing a cutoff δ_c and defining $C_{i,j} = \mathbb{1}[\Delta_{i,j} \leq \delta_c]$, then training

$$\min_{E,D} \ell(X, \sigma(CB_{E,D})). \quad (14)$$

C retains only opens fast enough to plausibly be observed within an active operational window and treats slower opens as non-opens during training — matching the censoring recipients are actually subject to during active learning. As $\delta_c \rightarrow \infty$, $C \rightarrow X$, recovering the original model.

3.7 Deep Autoencoder

$B_{E,D}$ in Eq. (2) is a single fixed $m \times m$ matrix, so the SAE can only capture linear recipient–recipient relationships. Since p_j and $f_j(t)$ depend only on the ability to query f_{SAE} , not on any explicit property of $B_{E,D}$, we can swap in a deeper, nonlinear network without changing either definition. We use a two-layer encoder/decoder with a ReLU nonlinearity, following the deep collaborative-filtering autoencoders of Kuchaiev and Ginsburg [2017], which were shown to generalize better than their shallow counterparts,

$$\text{Encoder: } \text{Lin}(m, 2d) \rightarrow \text{ReLU} \rightarrow \text{Lin}(2d, d), \quad \text{Decoder: } \text{Lin}(d, 2d) \rightarrow \text{ReLU} \rightarrow \text{Lin}(2d, m). \quad (15)$$

giving $f_{\text{DAE}}(x) = (\sigma \circ \text{Decoder} \circ \text{Encoder})(x)$, optionally with dropout [Srivastava et al., 2014] and bottleneck layer normalization. A deep network has no single weight matrix whose diagonal can be constrained to zero; we substitute a denoising objective [Vincent et al., 2008], randomly masking input entries and requiring reconstruction of the full unmasked X — as in collaborative denoising autoencoders for top-N recommendation [Wu et al., 2016] — forcing every prediction to depend on other recipients rather than a self-shortcut. Training becomes $\min_\theta \ell(X, f_{\text{DAE}}(X; \theta))$, and p_j , $f_j(t)$ retain their Sections 3.3–3.4 forms with f_{SAE} replaced by f_{DAE} .

4 Experiments

4.1 Methodology

For each variant we grid-search its hyperparameters and select the configuration maximizing validation-set Recall-AUC, $AUC = \int_0^1 \text{Recall}(\tau) d\tau$, where $\text{Recall}(\tau)$ is recall at sending fraction τ . We report test-set Recall@5/15/25/35% and Recall-AUC, mean $\pm$ std over repeated simulations. Before testing modifications we reproduced the original model on our pipeline; at the operating points reported by [Žid et al. \[2025\]](#) (25/50/75%) our reproduction matches the published point estimates exactly (0.923, 0.975, 0.989), with larger variance attributable to fewer simulation runs.

4.2 Results

Table 1: Test-set recall at four sending fractions and Recall-AUC (mean $\pm$ std) for each proposed refinement against the reproduced baseline. Hyperparameters were selected by grid search on the validation set.

Variant	Selected hyperparameters	Recall@5%	Recall@15%	Recall@25%	Recall@35%	AUC
Baseline	—	0.385 $\pm$ 0.076	0.821 $\pm$ 0.028	0.923 $\pm$ 0.014	0.958 $\pm$ 0.011	0.894 $\pm$ 0.012
Linear α -schedule	$l = 0.1, r = 0.05$	0.383 $\pm$ 0.073	0.823 $\pm$ 0.023	0.927 $\pm$ 0.012	0.960 $\pm$ 0.010	0.894 $\pm$ 0.011
Geometric α -schedule	$l = 0.3, r = 0.05$	0.377 $\pm$ 0.050	0.823 $\pm$ 0.017	0.927 $\pm$ 0.012	0.959 $\pm$ 0.010	0.894 $\pm$ 0.009
Confidence α -schedule	$\kappa = 0.003$	0.341 $\pm$ 0.031	0.818 $\pm$ 0.022	0.926 $\pm$ 0.014	0.960 $\pm$ 0.010	0.891 $\pm$ 0.009
Recency-weighted ϕ_j	$h = 0.3$	0.447 $\pm$ 0.041	0.837 $\pm$ 0.018	0.928 $\pm$ 0.014	0.960 $\pm$ 0.011	0.903 $\pm$ 0.009
Forward-pass $f_j(t)$	—	0.375 $\pm$ 0.049	0.841 $\pm$ 0.029	0.937 $\pm$ 0.018	0.966 $\pm$ 0.010	0.901 $\pm$ 0.011
Factored $s_j(t)$	$\alpha(t)$: conf., $\kappa = 0.005$; $\beta(t)$: geom., $l = 0.1, r = 0.05$	0.425 $\pm$ 0.026	0.821 $\pm$ 0.017	0.922 $\pm$ 0.013	0.959 $\pm$ 0.010	0.900 $\pm$ 0.008
Variance-based p_j	—	0.479 $\pm$ 0.034	0.837 $\pm$ 0.017	0.926 $\pm$ 0.013	0.960 $\pm$ 0.010	0.905 $\pm$ 0.009
Binary-to-TTO autoencoder	$h_\delta = 11520$	0.331 $\pm$ 0.051	0.719 $\pm$ 0.040	0.896 $\pm$ 0.015	0.955 $\pm$ 0.008	0.877 $\pm$ 0.011
TTO cutoff	$\delta_c = 720$	0.390 $\pm$ 0.038	0.824 $\pm$ 0.016	0.924 $\pm$ 0.012	0.960 $\pm$ 0.010	0.895 $\pm$ 0.008
Deep autoencoder	$d = 16$	0.404 $\pm$ 0.086	0.829 $\pm$ 0.052	0.931 $\pm$ 0.019	0.963 $\pm$ 0.011	0.900 $\pm$ 0.017

4.3 Discussion

Gains are largest at low sending fractions and shrink as τ grows, since baseline recall already exceeds 0.92 by $\tau = 25\%$. The **variance-based** p_j (0.479 at 5%), **recency-weighted** ϕ_j (0.447), and **factored** $s_j(t)$ (0.425) give the largest low-fraction gains and some of the highest AUC (0.905, 0.903, 0.900 vs. 0.894 baseline). The **forward-pass** $f_j(t)$ shows no gain at 5% but a consistent uplift from 15% onward, suggesting nonlinear aggregation only helps once significant number of recipients have opened. The α -scheduling variants are largely neutral, and the **Binary-to-TTO autoencoder** underperforms across all quantiles. **TTO cutoff** gives a small, consistent uplift at no interpretive cost, while the **deep autoencoder** improves recall modestly but with markedly higher variance, suggesting the added capacity is not yet well-regularized. All results test one modification at a time; combinations are untested.

5 Future Work

- **Combined α -scheduling.** A schedule jointly conditioned on send-progress μ and open-progress $\tilde{o}$ rather than either alone (Section 3.1); the main obstacle is the added hyperparameter count.
- **Plug-and-play composition under the factored score.** Systematically combine the primitive-level refinements — recency-weighted ϕ_j , forward-pass $f_j(t)$, variance-based p_j , TTO- or deep-autoencoder-derived Σ — inside $\pi_j(t)$ and $u_j(t)$ (Eq. (10)) rather than evaluating each in isolation, and search jointly rather than ablating one factor at a time, since the gains in Section 4.3 are not guaranteed to be additive.
- **Hyperbolic decay for Y .** $Y_{i,j} = \kappa_\delta / (\Delta_{i,j} + \kappa_\delta)$ gives a heavier tail than the exponential label of Eq. (12) and merits a direct comparison.
- **TTO-to-TTO autoencoder.** Replacing the SAE input with Y as well, $\min_{E,D} \ell(Y, \sigma(YB_{E,D}))$, exposes TTO on both sides but changes the role of e_j — an instantaneous open becomes a boundary point rather than a typical input — so p_j can no longer be queried at e_j and would need re-deriving, e.g. via a recipient- or population-level mean-TTO query, with unclear risk of double-counting against ϕ_j .

- **Asymmetric deep encoder/decoder.** The two-layer network of Section 3.7 uses matched encoder/decoder depth and width; decoupling them (e.g. a wider encoder with a narrower, more constrained decoder, or vice versa) may better balance representational capacity against the denoising-based regularization, and warrants exploration alongside a wider sweep over masking rate, dropout, and bottleneck width d to address the variance observed in Section 4.3.

References

- Jean-Yves Audibert, Rémi Munos, and Csaba Szepesvári. Exploration–exploitation tradeoff using variance estimates in multi-armed bandits. *Theoretical Computer Science*, 410(19):1876–1902, 2009. doi: 10.1016/j.tcs.2009.01.016.
- Peter Auer, Nicolò Cesa-Bianchi, and Paul Fischer. Finite-time analysis of the multiarmed bandit problem. *Machine Learning*, 47(2–3):235–256, 2002. doi: 10.1023/A:1013689704352.
- Yi Ding and Xue Li. Time weight collaborative filtering. In *Proceedings of the 14th ACM International Conference on Information and Knowledge Management (CIKM ’05)*, pages 485–492, Bremen, Germany, 2005. ACM. doi: 10.1145/1099554.1099689.
- S. Kirkpatrick, C. D. Gelatt, and M. P. Vecchi. Optimization by simulated annealing. *Science*, 220(4598):671–680, 1983. doi: 10.1126/science.220.4598.671.
- Yehuda Koren. Collaborative filtering with temporal dynamics. In *Proceedings of the 15th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD ’09)*, pages 447–456, Paris, France, 2009. ACM.
- Oleksii Kuchaiev and Boris Ginsburg. Training deep autoencoders for collaborative filtering. *arXiv preprint arXiv:1708.01715*, 2017.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Belle-mare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, Stig Petersen, Charles Beattie, Amir Sadik, Ioannis Antonoglou, Helen King, Dhharshan Kumaran, Daan Wierstra, Shane Legg, and Demis Hassabis. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, 2015. doi: 10.1038/nature14236.
- Suvash Sedhain, Aditya Krishna Menon, Scott Sanner, and Lexing Xie. AutoRec: Autoencoders meet collaborative filtering. In *Proceedings of the 24th International Conference on World Wide Web Companion (WWW ’15 Companion)*, pages 111–112. ACM, 2015. doi: 10.1145/2740908.2742726.
- Burr Settles. Active learning literature survey. Computer Sciences Technical Report 1648, University of Wisconsin–Madison, 2009.
- Moumita Sinha, Vishwa Vinay, and Harvineet Singh. Modeling time to open of emails with a latent state for user engagement level. In *Proceedings of the Eleventh ACM International Conference on Web Search and Data Mining (WSDM ’18)*, pages 531–539. ACM, 2018.
- Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Dropout: A simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15(56):1929–1958, 2014.

- Pascal Vincent, Hugo Larochelle, Yoshua Bengio, and Pierre-Antoine Manzagol. Extracting and composing robust features with denoising autoencoders. In *Proceedings of the 25th International Conference on Machine Learning (ICML '08)*, pages 1096–1103. ACM, 2008. doi: 10.1145/1390156.1390294.
- Čeněk Žid, Rodrigo Alves, and Pavel Kordík. Active recommendation for email outreach dynamics. In *Proceedings of the 34th ACM International Conference on Information and Knowledge Management (CIKM '25)*, pages 5540–5544, Seoul, Republic of Korea, 2025. ACM. doi: 10.1145/3746252.3760832.
- Yao Wu, Christopher DuBois, Alice X. Zheng, and Martin Ester. Collaborative denoising autoencoders for top-n recommender systems. In *Proceedings of the Ninth ACM International Conference on Web Search and Data Mining (WSDM '16)*, pages 153–162. ACM, 2016. doi: 10.1145/2835776.2835837.